

Social Data Science: Econometrics and Machine Learning

Week 2

Tree and kernel based methods

Social Data Science: Econometrics and Machine Learning

Week 2

~~Tree and kernel based methods~~

A tour of supervised machine learning (shallow learning)

Overview

1. Ensemble learning
2. Bagging, boosting
3. Decision trees
4. Random forests
5. K-nearest neighbor
6. Gradient boosting

Ensemble Learning

Ensemble Learning

Gist:

- Create and train many classification models
- Treat each model as a “voter”
- For each datapoint, classify it according to what models predicts it to be

Ensemble Learning

Gist:

- Create and train many classification models
- Treat each model as a “voter”
- For each datapoint, classify it according to what models predicts it to be

Pros:

1. Better generalization performance
2. Lowers overall error
3. Robust to overfitting

Cons:

1. Takes a little longer to train...

predictions of a single model can be highly sensitive to noise, but the average of many models is not

Bagging and boosting

Bagging and boosting

> Important concept: **bootstrapping**

Algorithm:

1. Given a list of length N , randomly select N elements *with* replacement

Bagging and boosting

> Important concept: **bootstrapping**

Algorithm:

1. Given a list of length N , randomly select N elements *with* replacement

In [10]:

```
1 import numpy as np
2
3 mylist = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
4 np.random.choice(mylist, size=len(mylist))
```

executed in 6ms, finished 07:11:28 2020-02-14

Out[10]: array([3, 8, 9, 3, 9, 2, 3, 0, 0, 3])

Bagging

> A machine learning strategy for ensemble learning

Algorithm:

Given a training set $X = x_1, \dots, x_n$ with responses $Y = y_1, \dots, y_n$, bagging repeatedly (B times) selects a **random sample with replacement** of the training set and fits trees to these samples:

For $b = 1, \dots, B$:

1. Sample, with replacement, n training examples from X, Y ; call these X_b, Y_b .
2. Train a classification or regression tree f_b on X_b, Y_b .

After training, predictions for unseen samples x' can be made by averaging the predictions from all the individual regression trees on x' :

$$\hat{f} = \frac{1}{B} \sum_{b=1}^B f_b(x')$$

Bagging

> A machine learning strategy for ensemble learning

Algorithm:

Given a training set $X = x_1, \dots, x_n$,
random sample with replacement

For $b = 1, \dots, B$:

- 1. Sample, with replacement
- 2. Train a classification or regression

After training, predictions for
individual regression trees on

$$\hat{f} = \frac{1}{B} \sum_{b=1}^B f_b(x')$$

Sample indices	Bagging round 1	Bagging round 2	...
1	2	7	...
2	2	3	...
3	1	2	...
4	3	1	...
5	7	1	...
6	2	7	...
7	4	7	...

C_1

C_2

C_m

repeatedly (B times) selects a
these samples:
se X_b, Y_b .
aging the predictions from all the

Bagging

> A machine learning strategy for ensemble learning

Algorithm:

Given a training set

random sample


For $b = 1, \dots, B$

1. Sample

2. Train a

After training

individual



$$\hat{f} = \frac{1}{B} \sum_{b=1}^B f_b(x')$$

es) selects a

ns from all the

Boosting

> A(nother) machine learning strategy for ensemble learning

Intuition: Many weak models that learn from each other's mistakes, combine into one strong model

Algorithm (most general):

1. Create a weight vector \mathbf{w} that encodes the *importance* of each training sample
2. For j out of m boosting iterations:
 - a. Train a weighted weak classifier $C_j = \text{train}(X, y, \mathbf{w})$
 - b. Predict class labels: $\hat{y} = \text{predict}(C_j, X)$
 - c. Update \mathbf{w} based on the errors that C_j makes (steps c-f in Raschka page 248)
3. To make predictions apply weighted voting, i.e. giving more prediction weight to less error prone classifiers

Boosting

> A(nother) machine learning strategy for ensemble learning

Intuition

Algorithm

- 1. Create
- 2. For
- a.
- b.
- c.
- 3. To
- classifiers

Sample indices	x	y	Weights	$\hat{y}(x \leq 3.0)?$	Correct?	Updated weights
1	1.0	1	0.1	1	Yes	0.072
2	2.0	1	0.1	1	Yes	0.072
3	3.0	1	0.1	1	Yes	0.072
4	4.0	-1	0.1	-1	Yes	0.072
5	5.0	-1	0.1	-1	Yes	0.072
6	6.0	-1	0.1	-1	Yes	0.072
7	7.0	1	0.1	-1	No	0.167
8	8.0	1	0.1	-1	No	0.167
9	9.0	1	0.1	-1	No	0.167
10	10.0	-1	0.1	-1	Yes	0.072

AdaBoost example from Raschka

less error prone

Boosting

> Mini exercise: discuss with your neighbor

Q1: What is the intuition behind weighting samples?

Q2: How is it practically done in AdaBoost (Raschka p. 248, c-f)

- c. Compute weighted error rate: $\varepsilon = \mathbf{w} \cdot (\hat{\mathbf{y}} \neq \mathbf{y})$.
- d. Compute coefficient: $\alpha_j = 0.5 \log \frac{1-\varepsilon}{\varepsilon}$.
- e. Update weights: $\mathbf{w} := \mathbf{w} \times \exp(-\alpha_j \times \hat{\mathbf{y}} \times \mathbf{y})$.
- f. Normalize weights to sum to 1: $\mathbf{w} := \mathbf{w} / \sum_i w_i$.

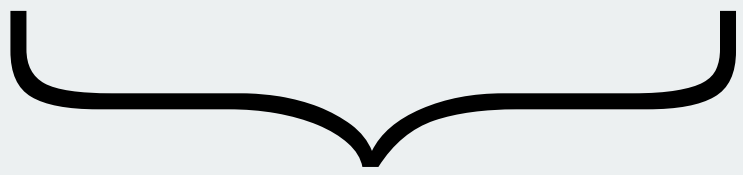
Q3: How are weights used during training?


$$C_j = \text{train}(X, y, \mathbf{w}).$$

Decision trees

Decision trees

Lays eggs	Cold blooded	Mammal
0	0	1
0	0	1
0	0	1
0	0	1
0	0	1
1	0	0
1	0	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0


 features


 target

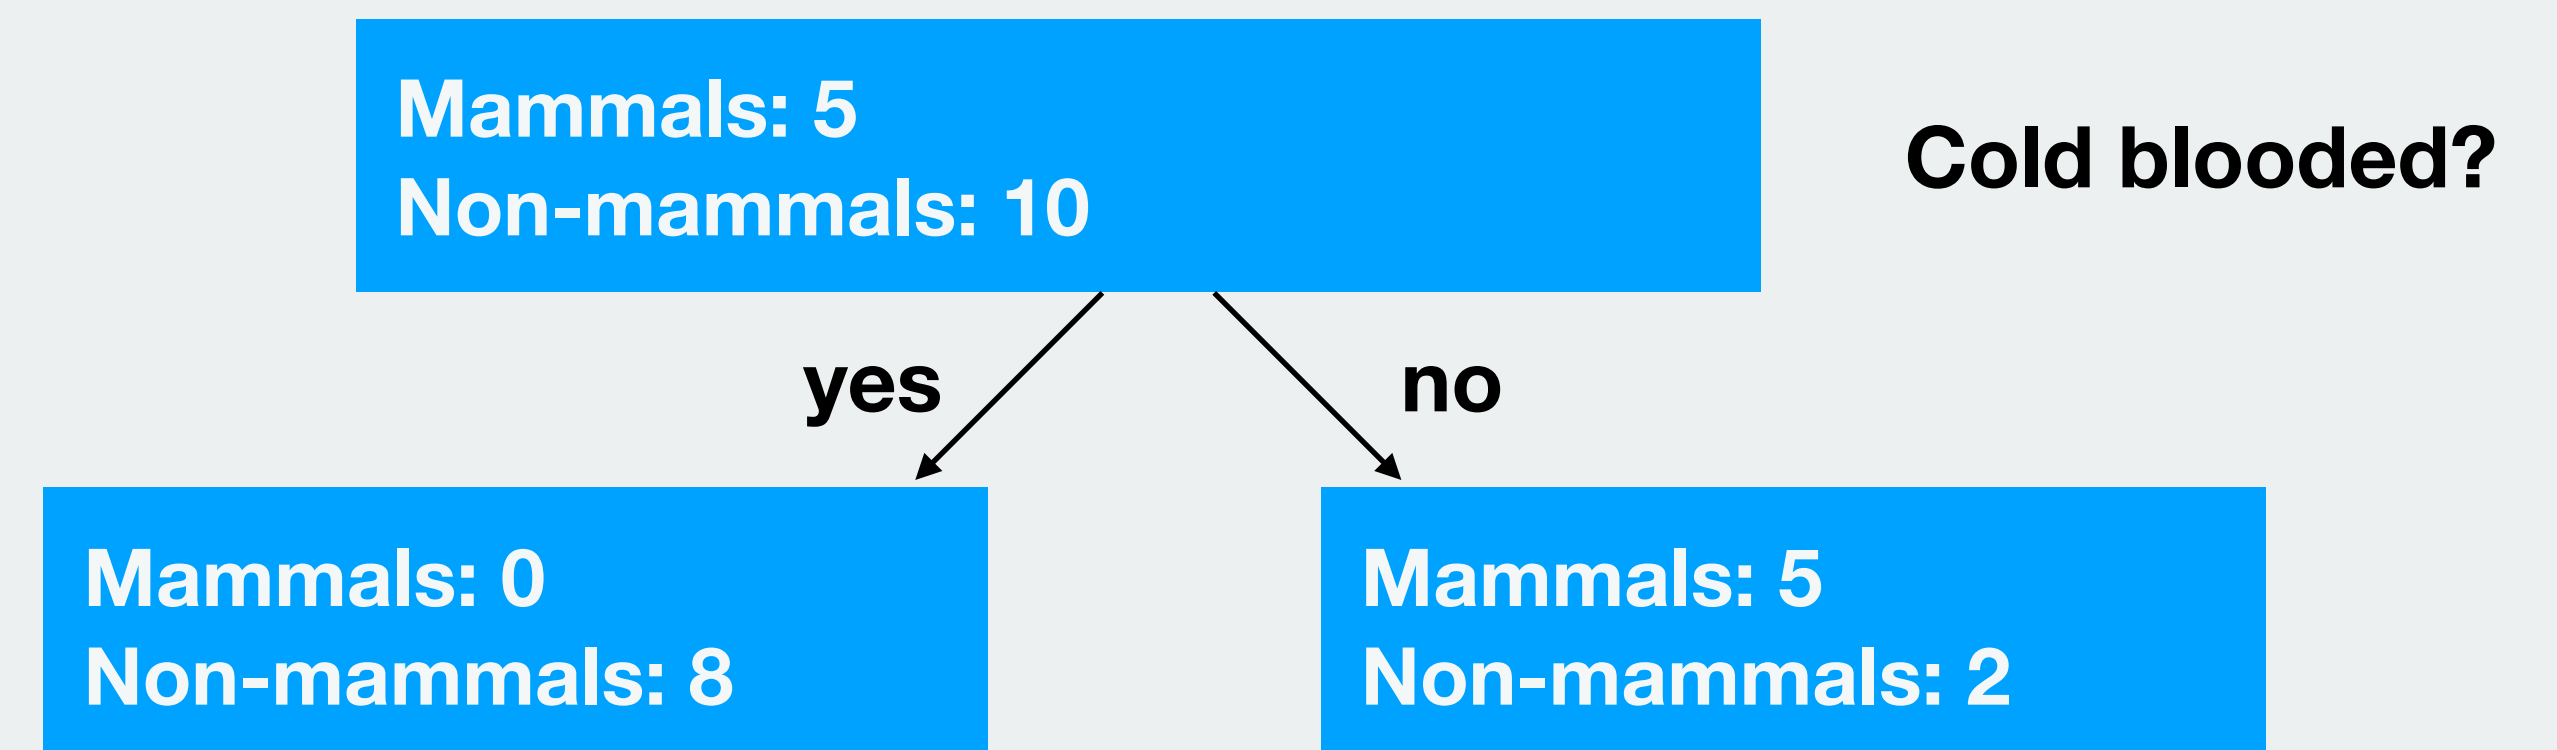
Mammals: 5
Non-mammals: 10

Decision trees


Lays eggs	Cold blooded	Mammal
0	0	1
0	0	1
0	0	1
0	0	1
0	0	1
1	0	0
1	0	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0
1	1	0

 features

 target



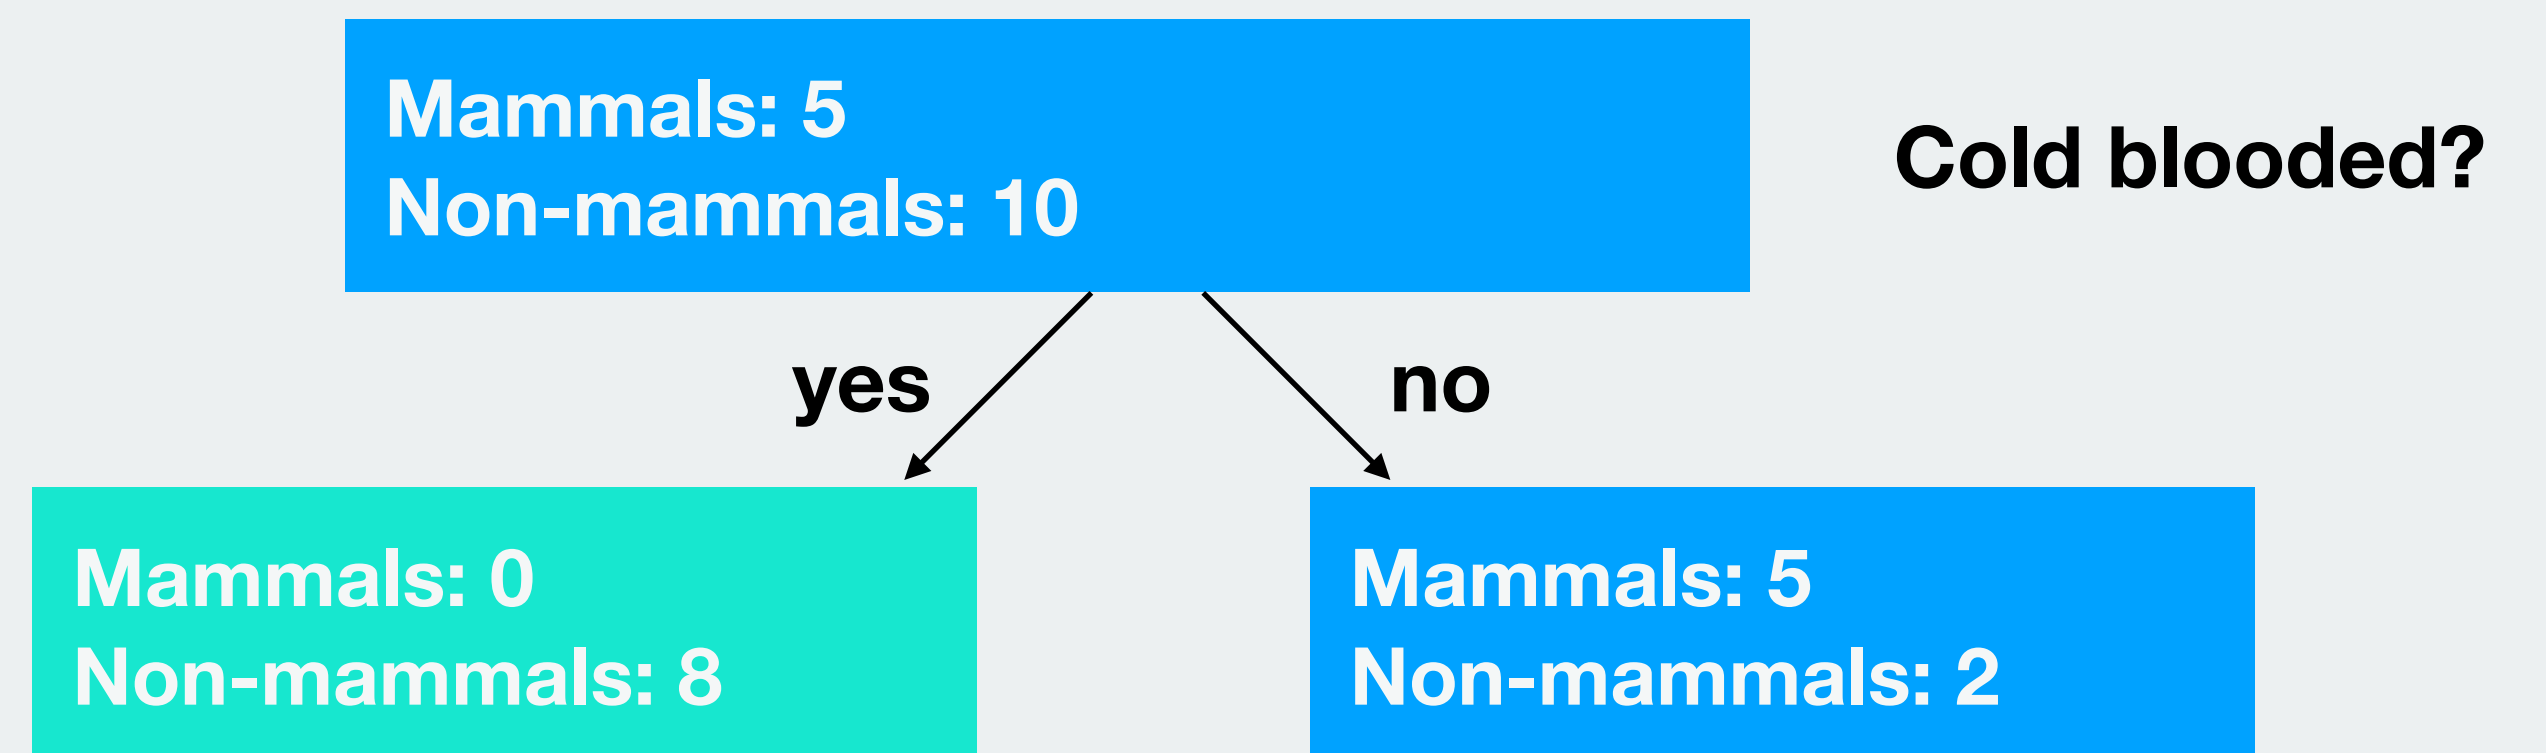
Decision trees



Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
	1	0	0
	1	0	0
	1	0	0
yes	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0

{ features
↑ target

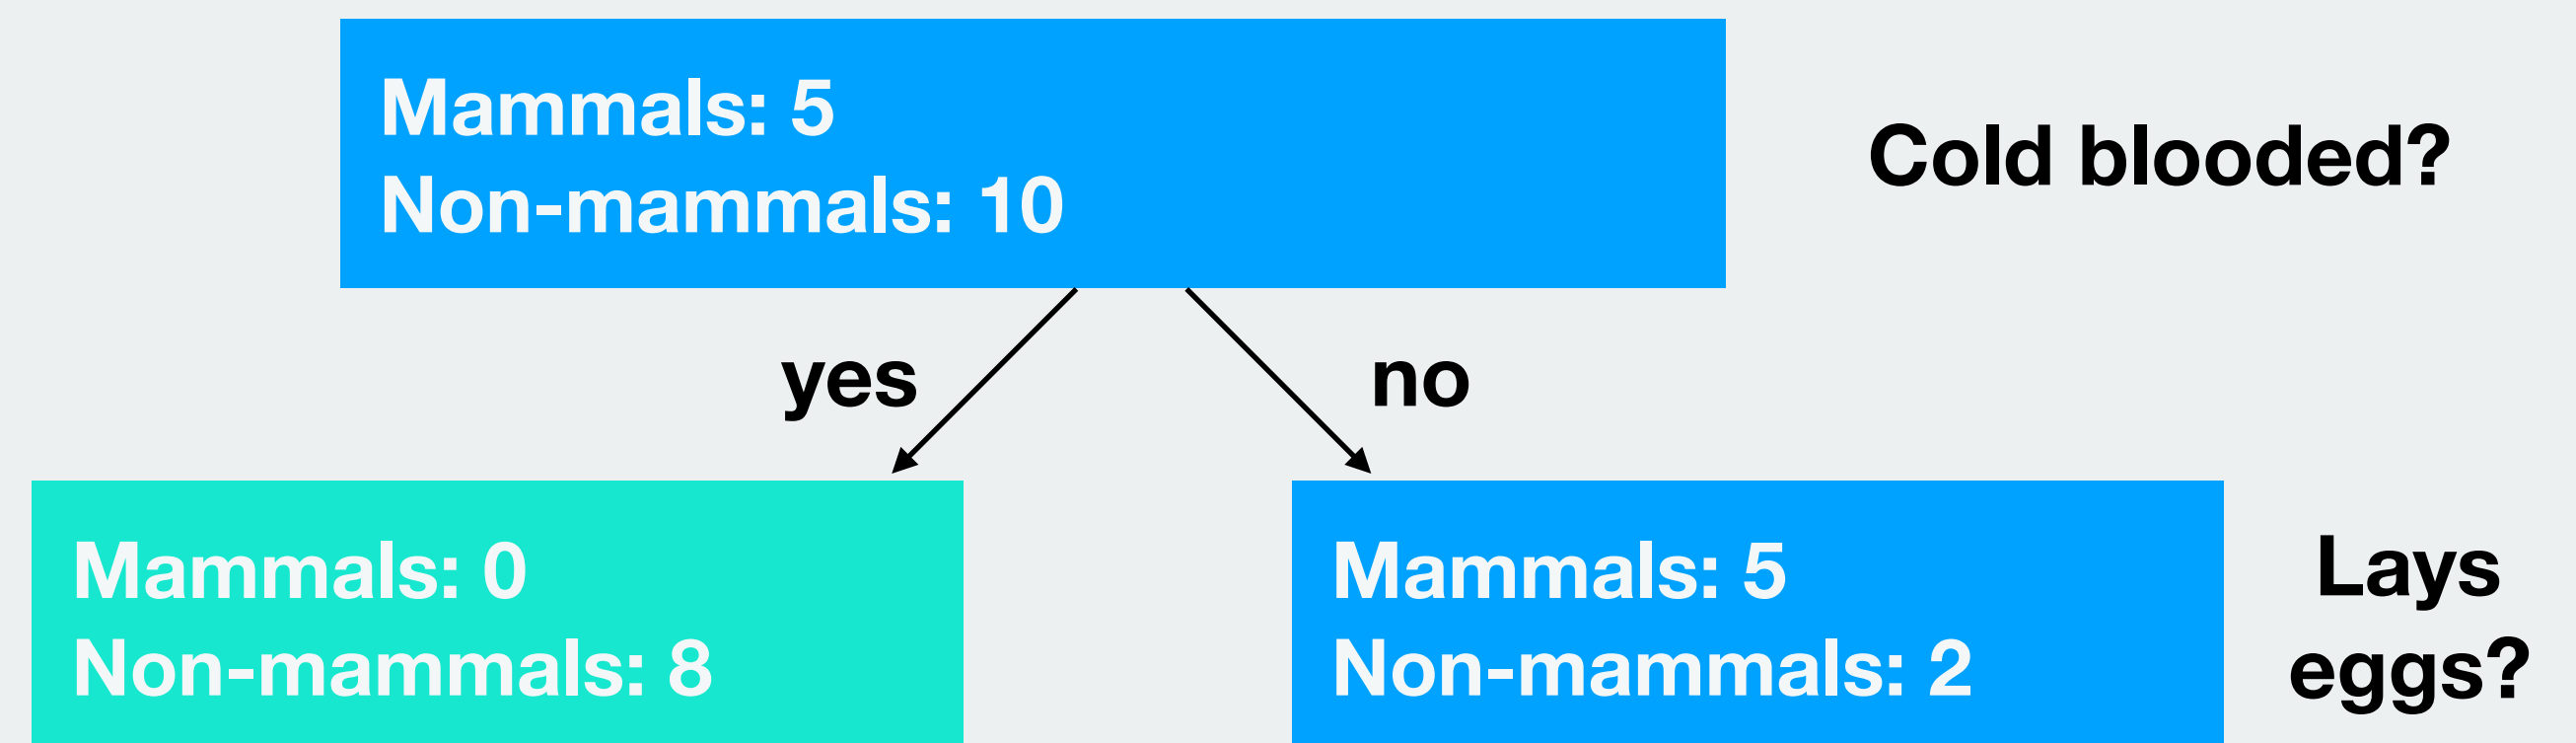


Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
	1	0	0
yes	1	0	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0

features


target



Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
yes	1	0	0
	1	0	0
yes	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0


{ features
↑ target



Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
yes	1	0	0
	1	0	0
yes	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0

features
target



Decision trees


Could we have asked better questions?

Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
yes	1	0	0
	1	0	0
yes	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0


features


target

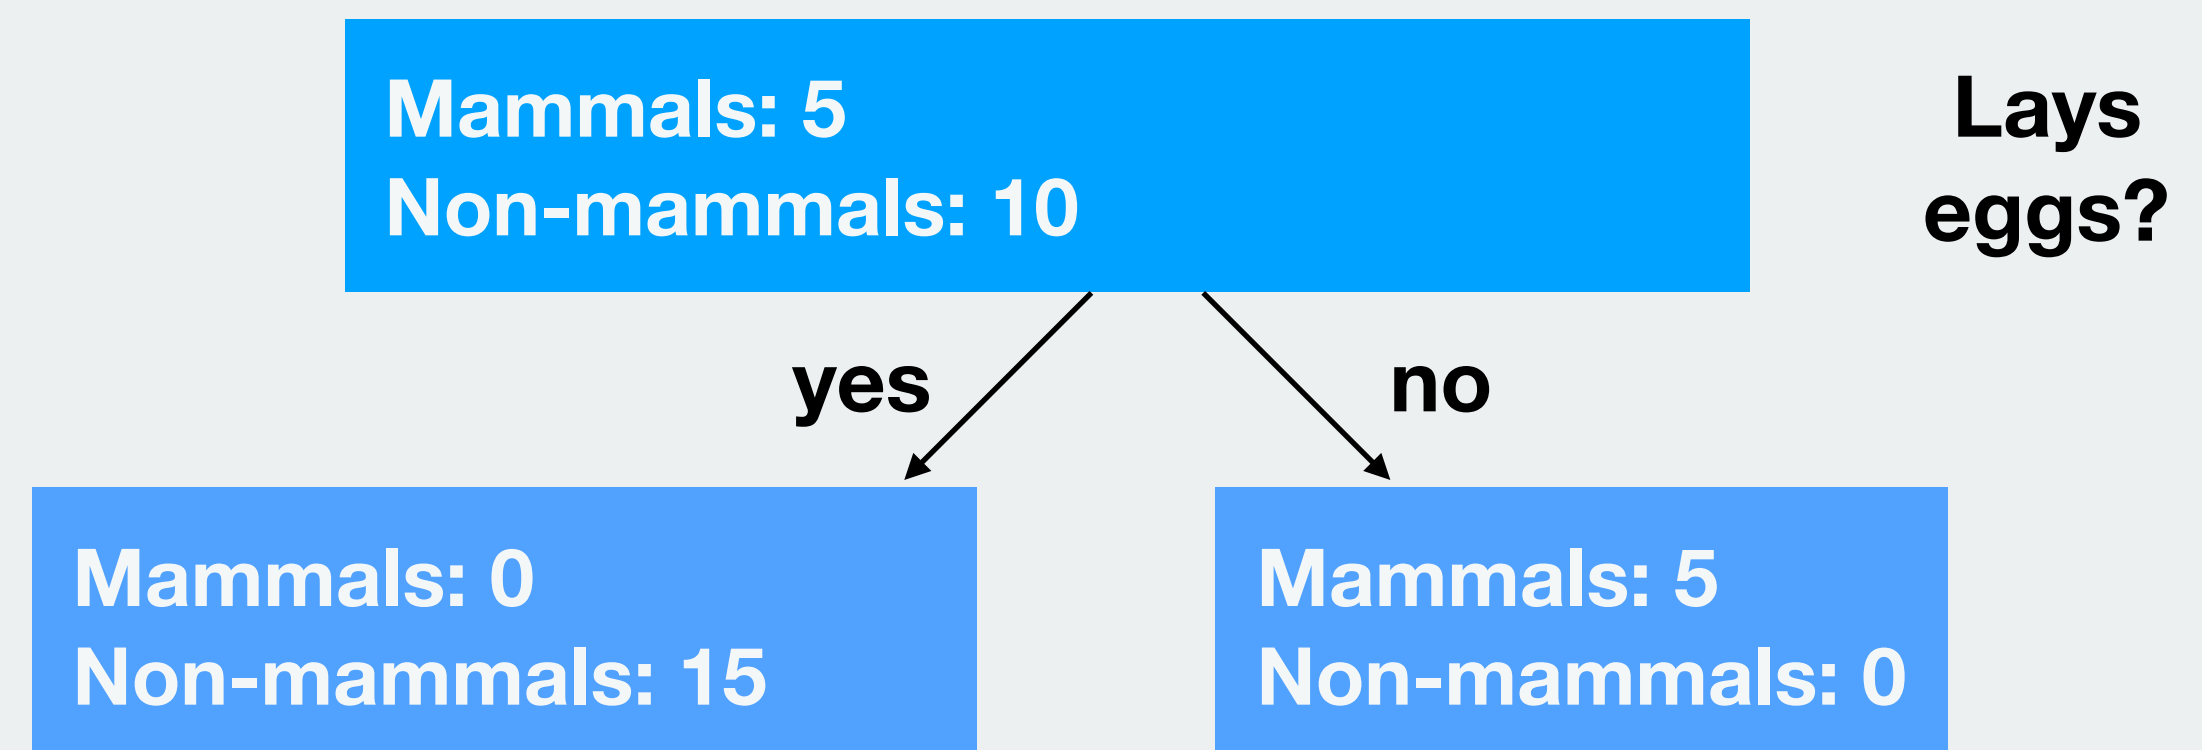


Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
yes	1	0	0
	1	0	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0

 features

 target

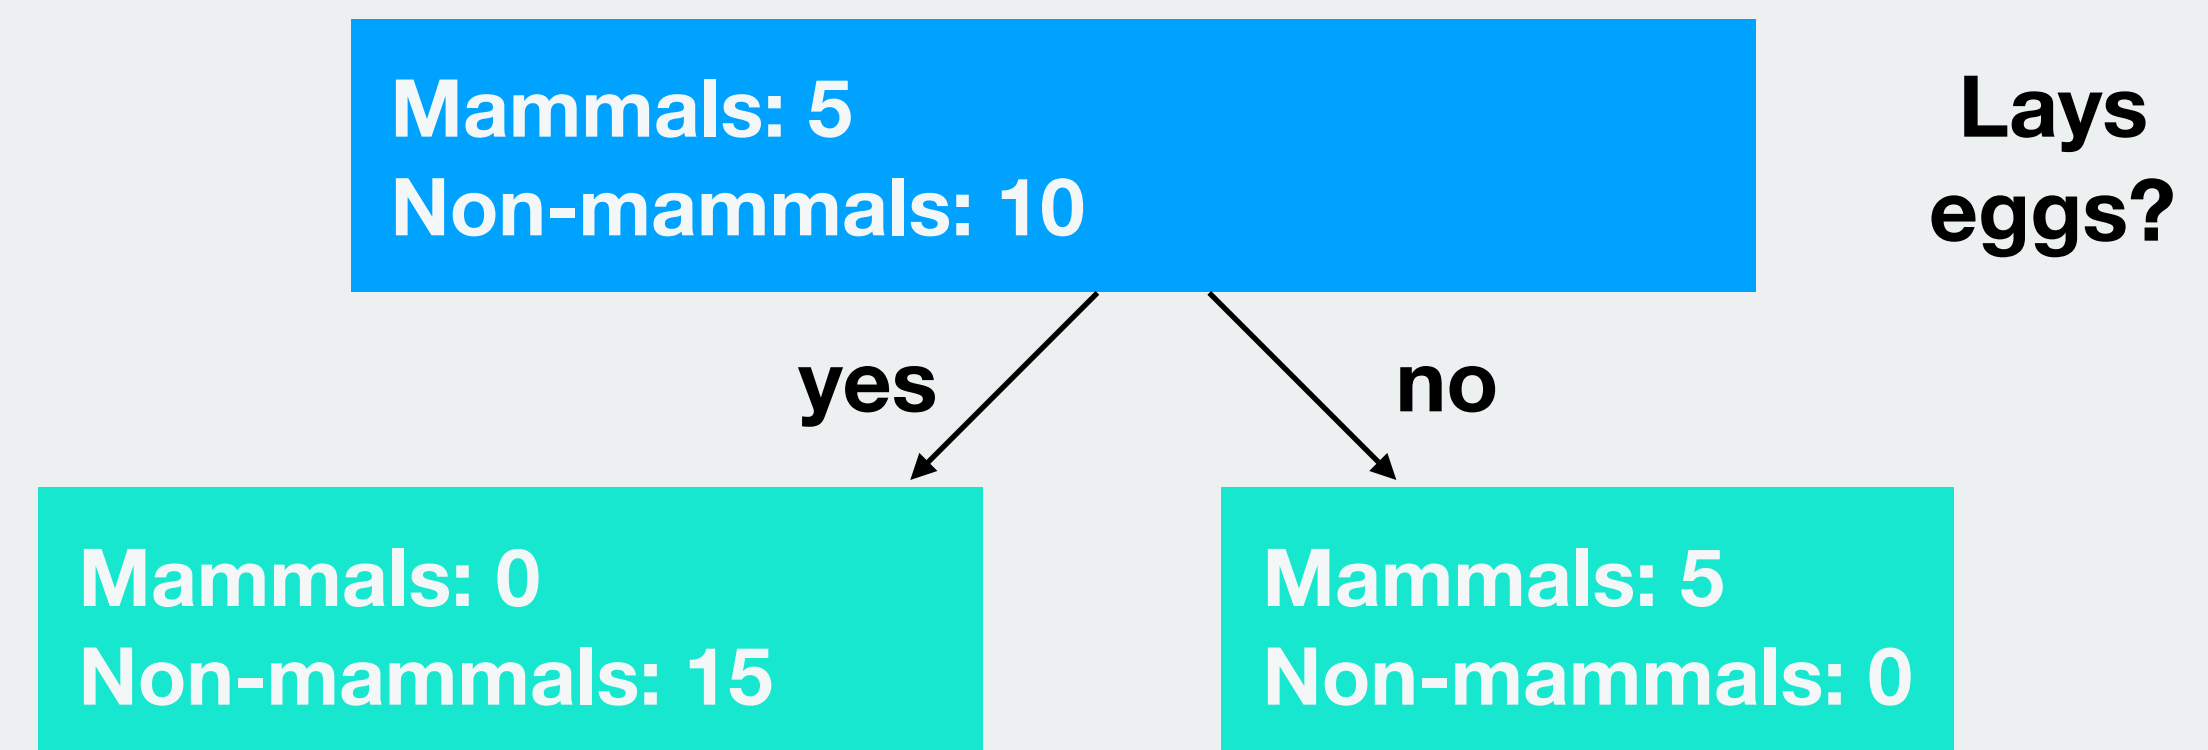


Decision trees

	Lays eggs	Cold blooded	Mammal
no	0	0	1
	0	0	1
	0	0	1
	0	0	1
	0	0	1
yes	1	0	0
	1	0	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0
	1	1	0

features

target



Decision trees

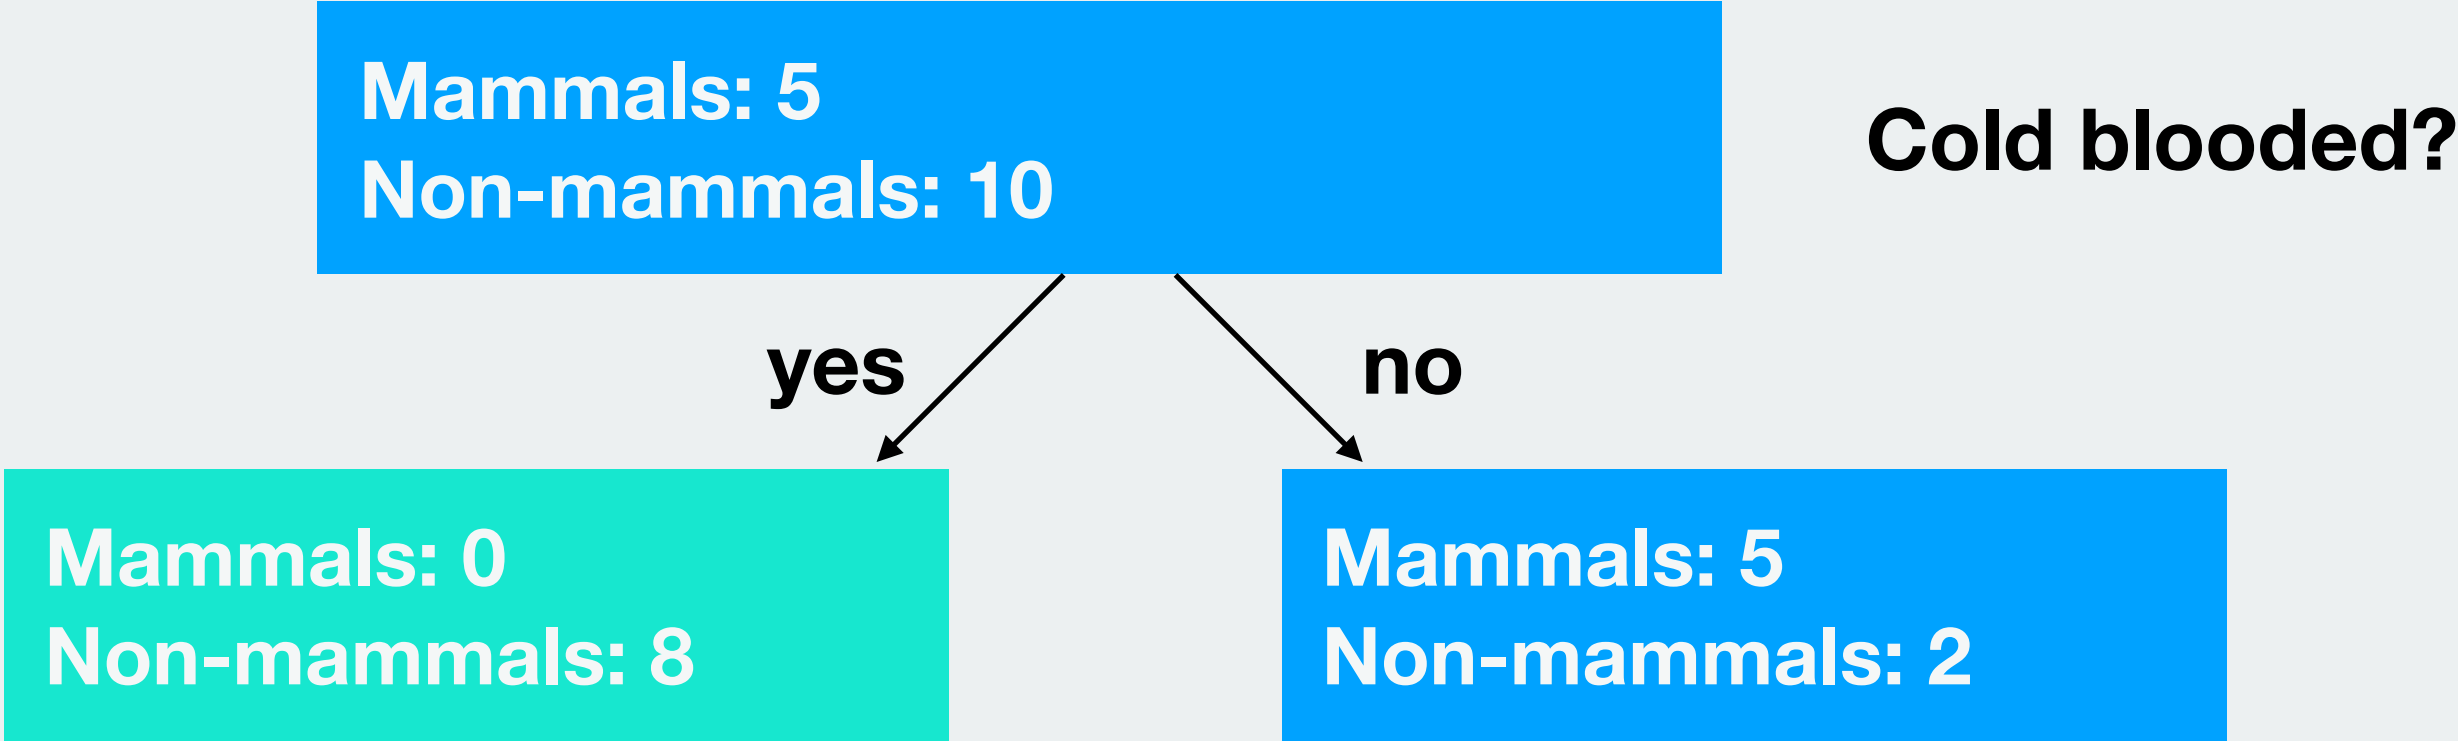
Can we somehow automate split selection?

Decision trees

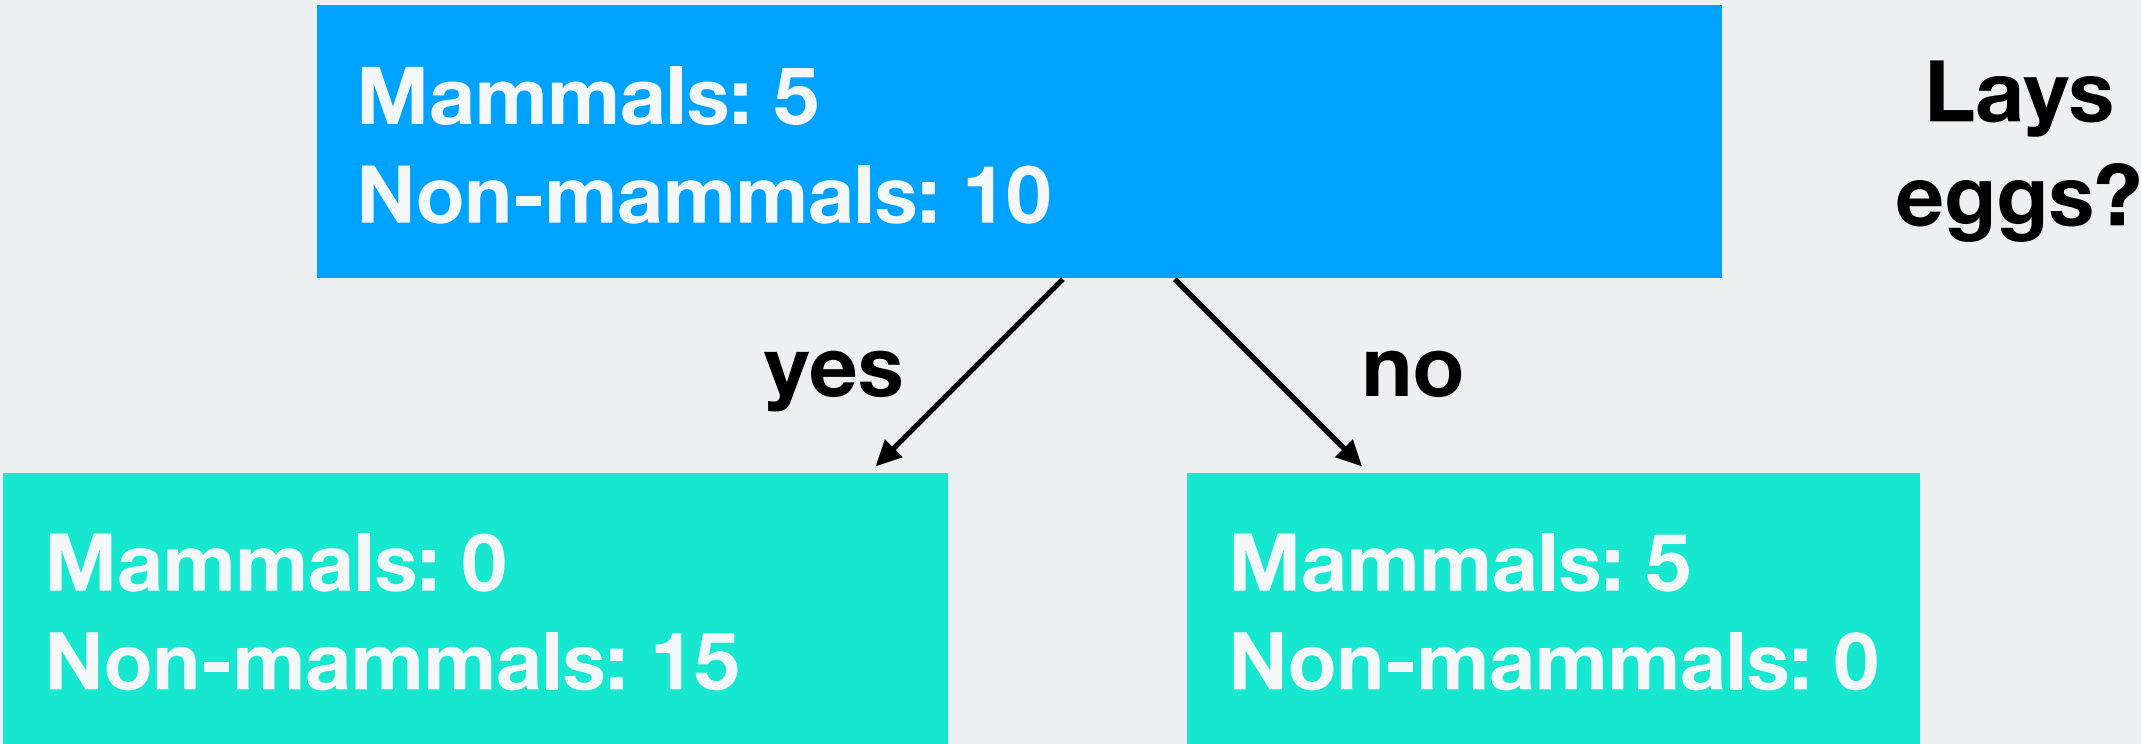
	Pclass1	Pclass2	Pclass3	Sexfemale	Sexmale	Embarkednan	EmbarkedC	EmbarkedQ	EmbarkedS	CabinFalse	CabinTrue	PassengerId	Age	SibSp	Parch	Fare	Survived
0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	1	22.0	1	0	7.2500	0
1	1.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	2	38.0	1	0	71.2833	1
2	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	3	26.0	0	0	7.9250	1
3	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	4	35.0	1	0	53.1000	1
4	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	5	35.0	0	0	8.0500	0
5	0.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	6	NaN	0	0	8.4583	0
6	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	7	54.0	0	0	51.8625	0
7	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	8	2.0	3	1	21.0750	0
8	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	9	27.0	0	2	11.1333	1
9	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	10	14.0	1	0	30.0708	1
...
881	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	882	33.0	0	0	7.8958	0
882	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	883	22.0	0	0	10.5167	0
883	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	884	28.0	0	0	10.5000	0
884	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	885	25.0	0	0	7.0500	0
885	0.0	0.0	1.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	886	39.0	0	5	29.1250	0
886	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	887	27.0	0	0	13.0000	0
887	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	888	19.0	0	0	30.0000	1
888	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	889	NaN	1	2	23.4500	0
889	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	890	26.0	0	0	30.0000	1
890	0.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	891	32.0	0	0	7.7500	0

Automatic split selection

Split 1:



Split 2:



Automatic split selection

$$\text{(Shannon)} \quad \textit{Entropy} = - \sum_i p(i) \log_2 p(i)$$

Input: Probability vector (a list of values between 0 and 1, which sums to 1)

Output: Entropy (a measure of how “spread out” the probability distribution is)

Automatic split selection

$$Entropy = - \sum_i p(i) \log_2 p(i)$$

Mammals: 0
Non-mammals: 8

$$p = [1, 0]$$

Automatic split selection

$$Entropy = - \sum_i p(i) \log_2 p(i)$$

Mammals: 0
Non-mammals: 8

$$p = [1, 0]$$

$$Entropy = - (1 \cdot \log_2(1) + 0 \cdot \log_2(0)) = 0$$

Automatic split selection

$$Entropy = - \sum_i p(i) \log_2 p(i)$$

Mammals: 0
Non-mammals: 8

$$p = [1, 0]$$

$$Entropy = - (1 \cdot \log_2(1) + 0 \cdot \log_2(0)) = 0$$


Mammals: 5
Non-mammals: 2

$$p = [2/7, 5/7]$$

$$Entropy = - (2/7 \cdot \log_2(2/7) + 5/7 \cdot \log_2(5/7)) = 0.86$$

Automatic split selection

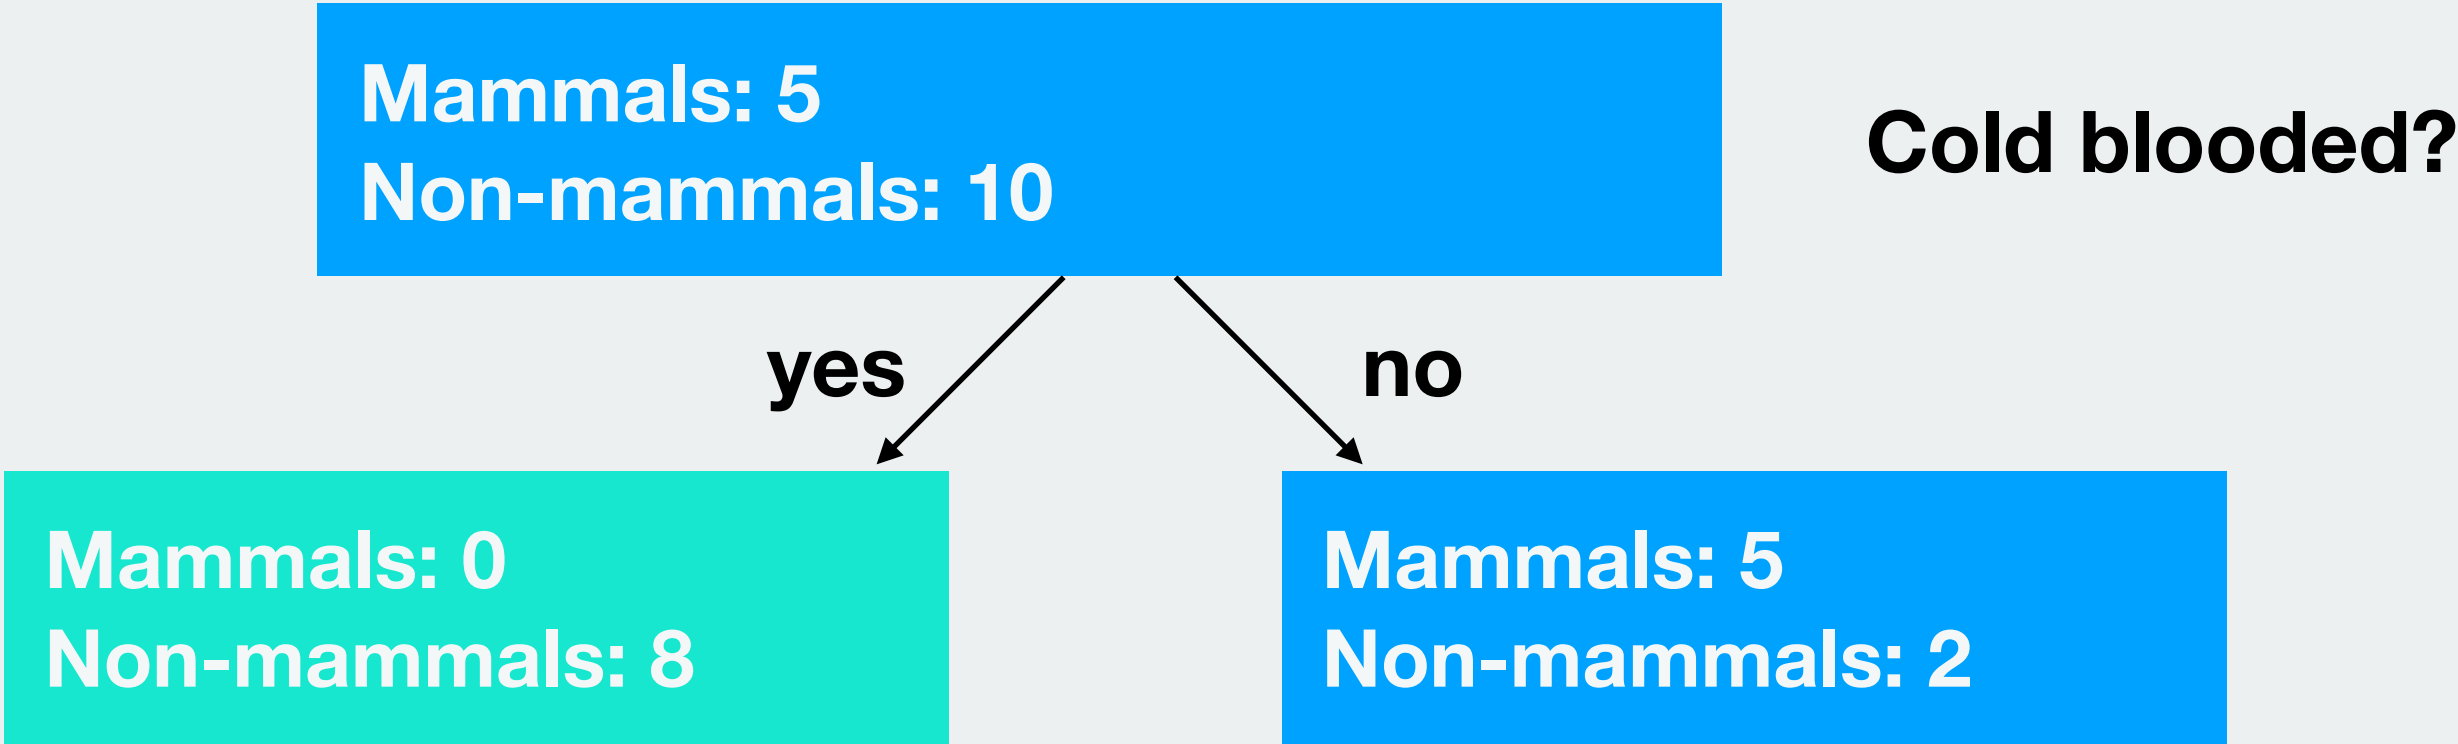
Split 1:



$$\text{split entropy} = 8 / 15 \cdot 0 + 7 / 15 \cdot 0.86 = 0.40$$

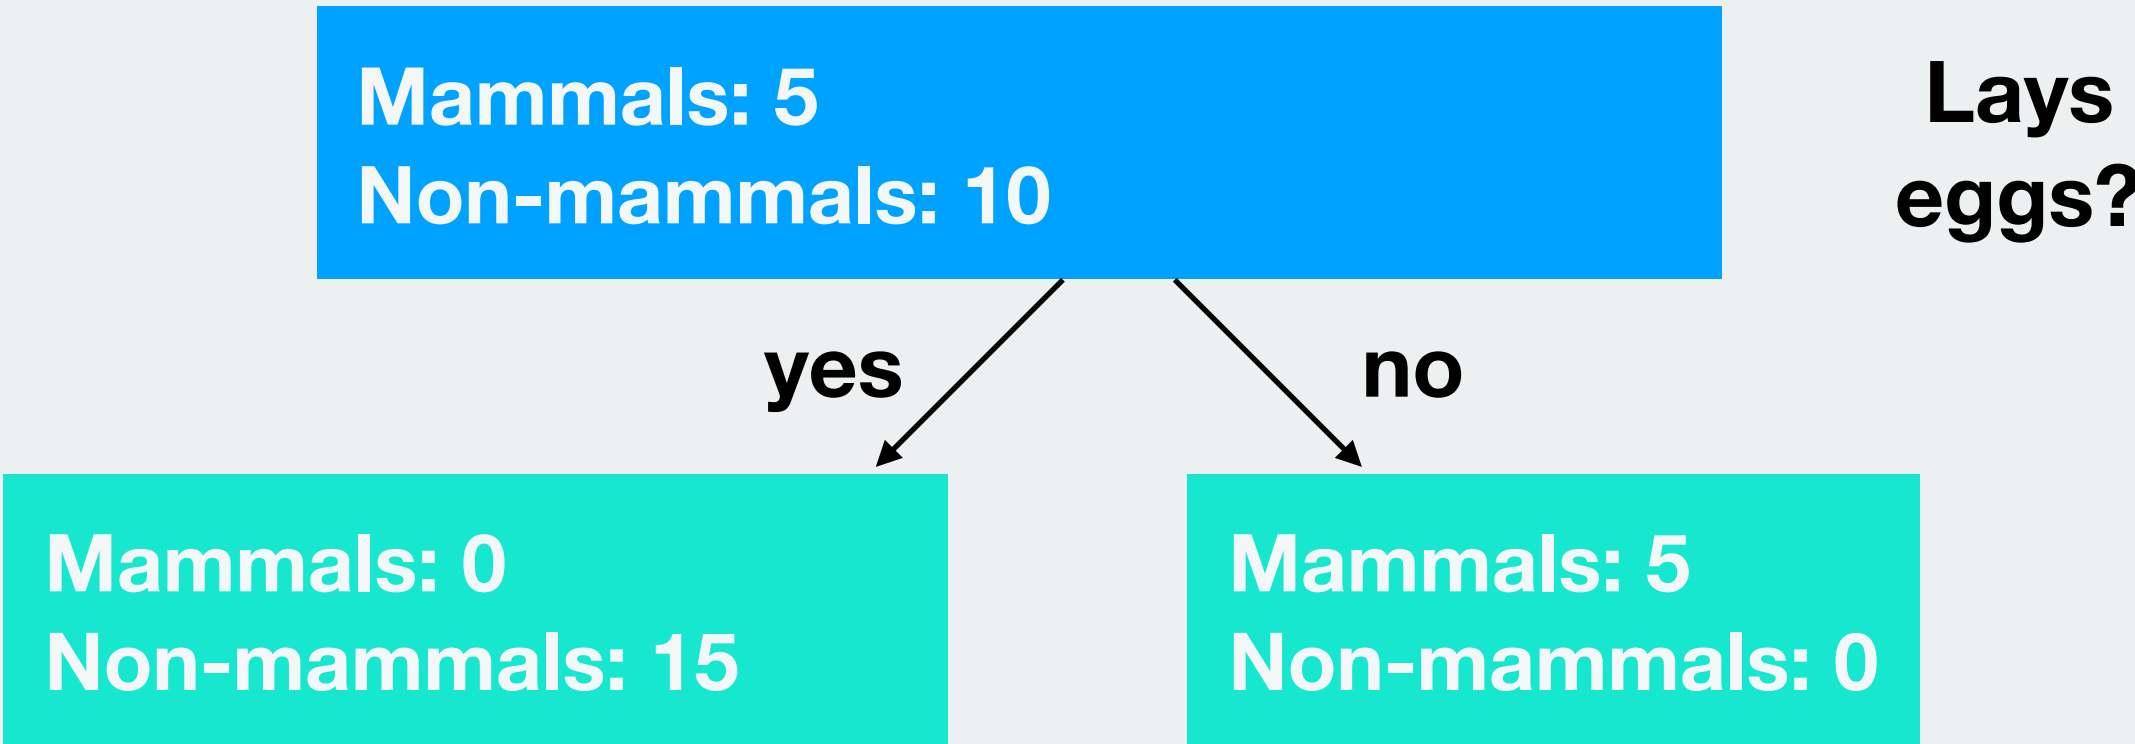
Automatic split selection

Split 1:



split entropy = $8 / 15 \cdot 0$ + $7 / 15 \cdot 0.86$ = 0.40

Split 2:



split entropy = $8 / 15 \cdot 0$ + $7 / 15 \cdot 0$ = 0

Random forests

Random forests

1. Draw a random **bootstrap** sample of size n (randomly choose n samples from the training set with replacement).

```
In [7]: 1 import numpy as np
        2
        3 mylist = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
        4 np.random.choice(mylist, size=10)
```

executed in 5ms, finished 21:36:59 2020-02-13

```
Out[7]: array([5, 8, 4, 9, 0, 8, 3, 9, 6, 0])
```

Random forests

1. Draw a random sample from the training set

	Pclass1	Pclass2	Pclass3	Sexfemale	Sexmale	Embarkednan	EmbarkedC	EmbarkedQ	EmbarkedS	CabinFalse	CabinTrue	PassengerId	Age	SibSp	Parch	Fare	Survived
0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	1	22.0	1	0	7.2500	0
1	1.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	2	38.0	1	0	71.2833	1
2	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	3	26.0	0	0	7.9250	1
3	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	4	35.0	1	0	53.1000	1
4	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	5	35.0	0	0	8.0500	0
5	0.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	6	NaN	0	0	8.4583	0
6	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	7	54.0	0	0	51.8625	0
7	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	8	2.0	3	1	21.0750	0
8	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	9	27.0	0	2	11.1333	1
9	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	10	14.0	1	0	30.0708	1
...
881	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	882	33.0	0	0	7.8958	0
882	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	883	22.0	0	0	10.5167	0
883	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	884	28.0	0	0	10.5000	0
884	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	885	25.0	0	0	7.0500	0
885	0.0	0.0	1.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	886	39.0	0	5	29.1250	0
886	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	1.0	0.0	887	27.0	0	0	13.0000	0
887	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	888	19.0	0	0	30.0000	1
888	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	889	NaN	1	2	23.4500	0
889	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	1.0	890	26.0	0	0	30.0000	1
890	0.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	891	32.0	0	0	7.7500	0

from the

Random forests


1. Draw a random **bootstrap** sample of size n (randomly choose n samples from the training set with replacement).
2. Grow a decision tree from the bootstrap sample. At each node:
 - Randomly select d features without replacement.
 - Split the node using the feature that provides the best split according to the objective function, for instance, maximizing the information gain.

Random forests

1. Draw a random bootstrap sample from the training set with replacement

2. Grow a decision tree

- Randomly select features to split on
- Split the node based on a splitting function, for instance, the Gini index



3. Draw samples from the

4. Aggregate the results according to the objective

Random forests

1. Draw a random **bootstrap** sample of size n (randomly choose n samples from the training set with replacement).
2. Grow a decision tree from the bootstrap sample. At each node:
 - Randomly select d features without replacement.
 - Split the node using the feature that provides the best split according to the objective function, for instance, maximizing the information gain.
3. Repeat the steps 1-2 k times.
4. Aggregate the prediction by each tree to assign the class label by **majority vote**.

Random forests

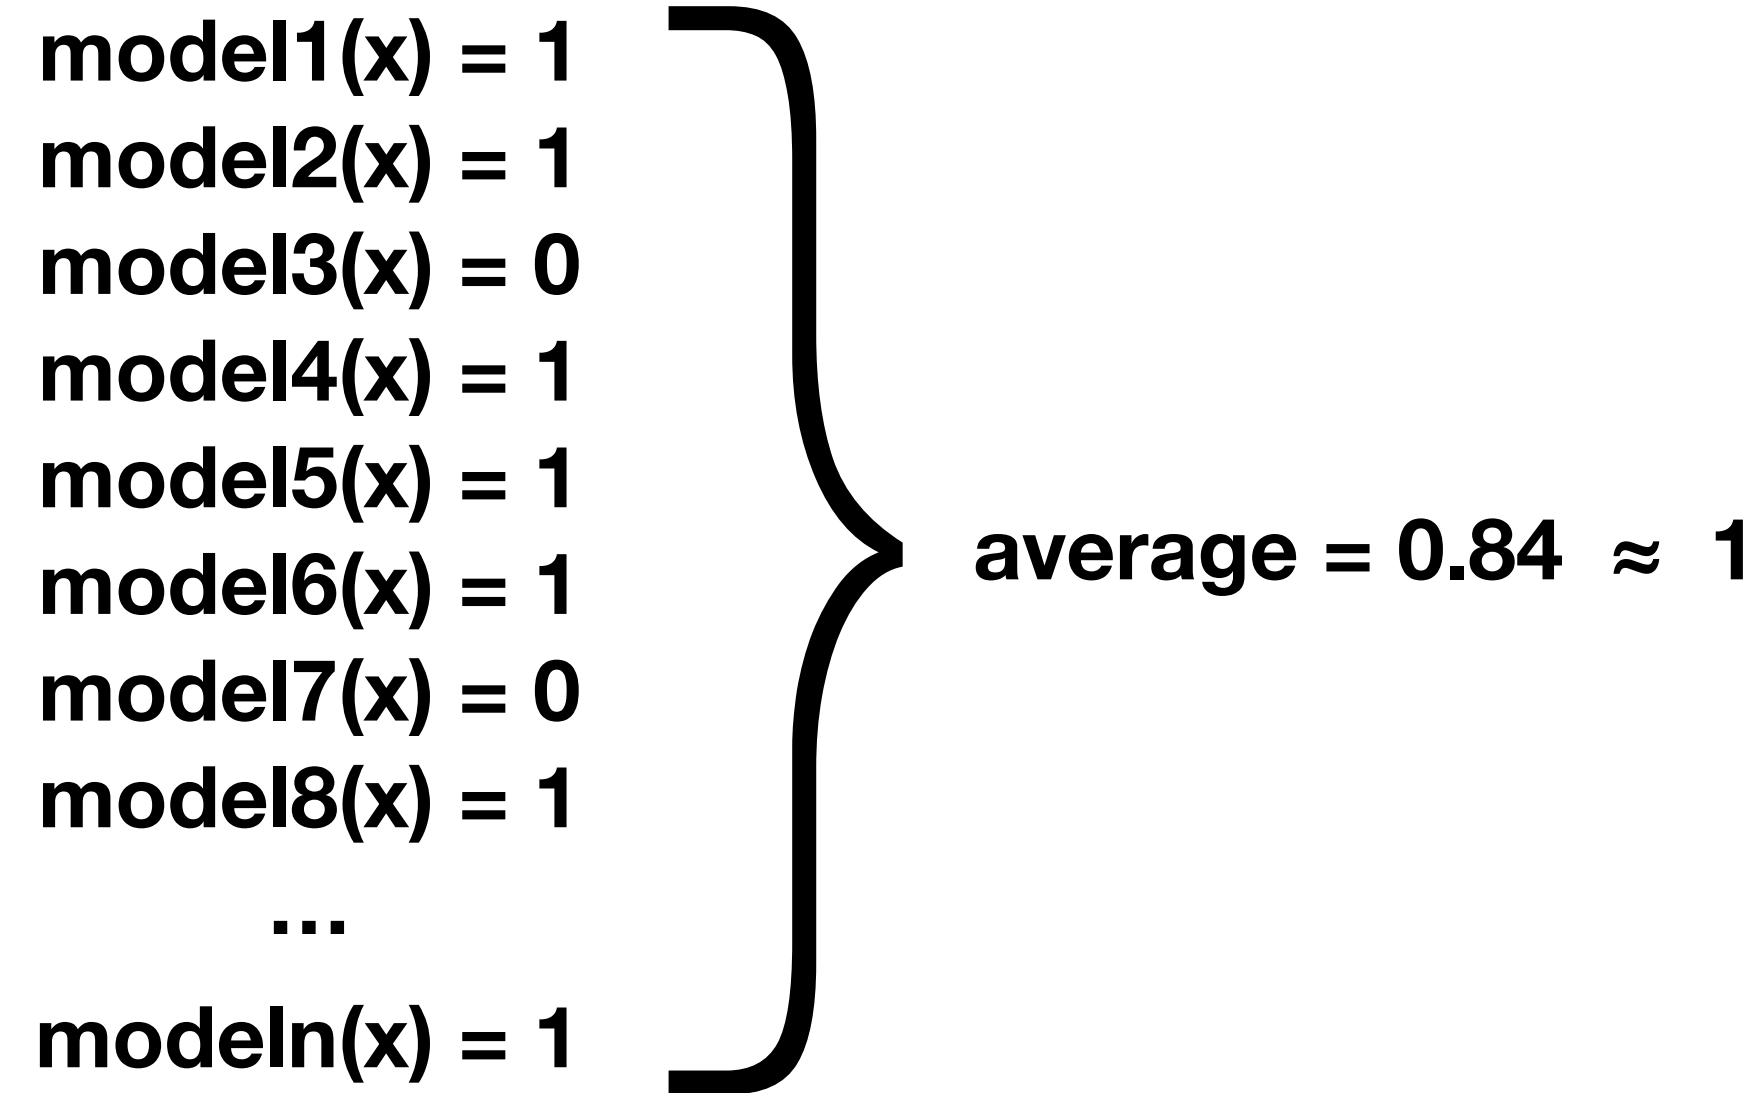
1. Draw a random **bootstrap** sample of n samples from the training set with replacement

2. Grow a decision tree from the bootstrap sample

- Randomly select d features
- Split the node using a splitting function, for instance

3. Repeat the steps 1-2 k times

4. Aggregate the prediction by each tree to assign the class label by **majority vote**.


$$\begin{array}{l} \text{model1}(x) = 1 \\ \text{model2}(x) = 1 \\ \text{model3}(x) = 0 \\ \text{model4}(x) = 1 \\ \text{model5}(x) = 1 \\ \text{model6}(x) = 1 \\ \text{model7}(x) = 0 \\ \text{model8}(x) = 1 \\ \dots \\ \text{modeln}(x) = 1 \end{array} \quad \left. \vphantom{\begin{array}{l} \text{model1}(x) = 1 \\ \text{model2}(x) = 1 \\ \text{model3}(x) = 0 \\ \text{model4}(x) = 1 \\ \text{model5}(x) = 1 \\ \text{model6}(x) = 1 \\ \text{model7}(x) = 0 \\ \text{model8}(x) = 1 \\ \dots \\ \text{modeln}(x) = 1 \end{array}} \right\} \text{average} = 0.84 \approx 1$$

n samples from the

according to the objective

Random forests

> Quick word on feature importance

- When fitting trees, we estimate the information gain for every feature at each split
- We can summarize the *importance* of a feature as its relative amount of information gain delivered during classification

```
In [18]: 1 from sklearn.ensemble import RandomForestClassifier
          2
          3 model = RandomForestClassifier()
          4 model.fit(X, y)
          5 model.feature_importances_

executed in 50ms, finished 09:38:37 2020-02-14

Out[18]: array([0.15166831, 0.1353509 , 0.25520054, 0.27534356, 0.03120302,
                0.01501129, 0.01015842, 0.00808299, 0.01251548, 0.00859388,
                0.00075693, 0.          , 0.          , 0.01614362, 0.02831235,
                0.02773086, 0.02392785])
```

K-nearest neighbors

K-nearest neighbors


Algorithm:

1. Choose the number of k and a distance metric.
2. Find the k -nearest neighbors of the sample that we want to classify.
3. Assign the class label by majority vote.

K-nearest neighbors

Algorithm:

1. Choose the number
2. Find the k-nearest
3. Assign the class label





classify.

K-nearest neighbors

Algo

1. Choose
2. Find
3. Assign



K-nearest neighbors

Pros:

Cons:

K-nearest neighbors

Pros:

1. No training needed!
2. Immediately adapts as we add more training data

Cons:

1. Search space grows linearly with amount of data
2. Choosing K requires hyperparameter tuning
3. Prone to overfitting due to ‘Curse of dimensionality’
4. For most distance metrics, data must be standardized

Gradient boosting

Gradient boosting

Algorithm (informally):

1. Create a very weak model F_0 that just predicts the average class, \bar{y}
2. Estimate the vector of residuals $y - F_0(X)$
3. Fit another weak model F_1 to predict the residuals $y - F_0(X)$. A perfect F_1 would imply $F_{0+1} = F_0(X) + F_1(X) = y$
4. But F_1 is weak, and highly imperfect. Therefore, continue adding more functions like this for many more iterations

Gradient boosting

Pros:

1. Very good performance with classification and regression trees
2. Easily extensible with arbitrary cost functions
3. State-of-the-art performance on *shallow* learning problems
4. The XGBoost library makes it easy to use

Cons:

1. Training is not super fast
2. Overfitting must be controlled with regularization

Further learning

Check out:

1. Scikit-learn user guide for a practical overview of methods
https://scikit-learn.org/stable/user_guide.html
2. The XGBoost library for gradient boosting
<https://xgboost.readthedocs.io/en/latest/>
- 3.